

PRITI PRADEEP GHOSH

Boston, MA | (617) 309-7849 | ghoshpriti30@gmail.com | [Linkedin](#) | [Github](#) | [Portfolio](#)

SUMMARY

Data Engineer with 3+ years of experience in designing and automating large-scale ETL/ELT architectures across AWS and Azure environments. Expert in PySpark and SQL optimization for multi-terabyte datasets, with a proven record of reducing manual debugging. Adept at building data lakes and monitoring solutions to drive data-driven business decisions.

EDUCATION

Northeastern University

Master of Science, Computer Software Engineering

Jan 2024 - Apr 2026

Boston, MA

University of Mumbai

Bachelor of Engineering, Electronics and Telecommunication

Jun 2017 - May 2021

Mumbai, India

TECHNICAL SKILLS

- **Programming & Frameworks:** Python, Pandas, PySpark, SQL, T-SQL, PostgreSQL, MySQL, Hadoop, RDBMS, Java, JSON
- **Data Engineering & Orchestration:** ETL/ELT/ETLT Pipelines, AWS Step Functions, EventBridge, Azure Data Factory
- **Cloud Architecture (AWS):** Glue, Jobs, Crawlers, DataBrew, S3, Data Lake, Redshift, Lambda, Athena, DynamoDB
- **Modern Data Stack & Warehousing:** Snowflake, Data Lakehouse Architecture, Data Modeling, Parquet/Avro
- **DevOps & DataOps:** CI/CD Pipelines, Azure DevOps, IAC, CloudFormation, Git, Docker, CloudWatch
- **Governance & Security:** IAM, Fine-grained permissions, Palantir Data Governance, Schema Validation, DQM
- **Visualization & Analytics:** Power BI, Tableau, Alteryx, Talend, MS Excel
- **Coursework:** Agile Software Development, Data Architecture and Business Intelligence, Data Science, Prompt Engineering
- **Certifications:** AWS Cloud Data Engineer, AZ-104/900 Microsoft Azure Fundamentals & Associate, MS Power BI, Snowflake

WORK EXPERIENCE

Ada IQ

Data Engineering Co-op

Jan 2025 - May 2025

Boston, USA

- Architected high-throughput ETL pipelines using AWS Glue and Step Functions to ingest 2TB+ of daily telemetry data, directly enabling the Data Science team to deliver real-time predictive modeling for business solutions.
- Engineered a self-healing Data Quality Framework utilizing AWS DataBrew and Athena to automate schema validation and anomaly detection, ensuring 99.9% data reliability for executive-level reporting.
- Reduced cloud infrastructure overhead by 20% by implementing S3 Intelligent-Tiering and optimizing partitioning in Athena and Redshift, improving query efficiency, minimizing data scan cost, enhancing performance of analytics workload.
- Spearheaded an Automated DevOps monitoring suite using CloudWatch, Lambda, and EventBridge to automate pipeline failure alerts via Slack, reducing MTTR (Mean Time to Resolution) by 75%.
- Partnered with client and stakeholders in a fast-paced startup environment, operating within Agile sprints to strengthen data governance, ensuring audit readiness while maintaining seamless access for business users and external partners.

Accenture

Data Engineer Analyst

Jun 2021 - Jan 2024

Mumbai, India

- Architected and deployed enterprise-grade CI/CD workflows via Azure DevOps and CloudFormation implementing IaC, achieving a 90% reduction in deployment time and accelerating the data development lifecycle by 25%.
- Spearheaded the migration of legacy SQL environments to AWS Redshift; re-architected schemas for Parquet data sources for Columnar Storage, resulting in a 5x improvement in query performance for global BI analytics.
- Engineered a cross-platform monitoring dashboard using Azure Monitor and Log Analytics to track ADF pipelines; slashed critical incident response times by 30% and ensured 24/7 business continuity.
- Developed an automated Python-based Data Governance Framework for Palantir datasets, implementing validation checks that ensured 100% regulatory compliance and eliminated data quality bottlenecks.
- Engaged directly with enterprise clients and cross-functional stakeholders in an Agile delivery model to gather requirements, manage technical trade-offs, and deliver scalable data solutions aligned with strategic business objectives.

ACADEMIC PROJECTS

Vehicle Collision Analysis: End to End Data Engineering Project

- Aggregated 3M+ records using Alteryx for profiling and Talend/ADF for transformation; performed advanced data cleaning and feature engineering using Python (Pandas)
- Architected a Medallion data model (Bronze, Silver, Gold) in Snowflake; implemented SCD Type 2 using surrogate keys and effective date tracking to preserve historical changes and support auditability.
- Engineered optimized views in Snowflake to power KPI-driven dashboards in Power BI and Tableau, delivering actionable insights on collision trends while reducing query latency via Parquet-based storage.